Naman Parmar

+91 9537112498 | Anand, India | parmarnaman19@gmail.com
linkedin.com/in/namanparmar | github.com/Naman0807 | huggingface.co/Naman0807 | naman0807.github.io/myself

Summary

AI-first developer shipping a production LLM tool-calling agent (Azure OpenAI, GPT-6-Luna) in a HIPAA-regulated healthcare platform. Trained Gujarati GPT-2 and LLaMA2 models from scratch. Seeking AI engineer roles.

Experience

AI-First Developer June 2026 - Present
GroovyWeb *Anand, India*

- Built the AI assistant for a HIPAA-regulated US healthcare patient platform: a **tool-using agent on Azure OpenAI (GPT-6-Luna)** with 7 typed function tools answering from the patient's own visits, medications, consent status, rewards, and saved studies.
- Designed the **multi-turn agent loop with memory and a turn cap** bounding per-conversation cost, behind a provider interface: deterministic stub in CI, Azure in production, zero network calls in tests.
- Built **prompt guardrails**: intent classification plus pre- and post-filters refusing medical advice, off-domain requests, and compensation as a reason to attend a visit, on both sides of the model call.
- Wrote a **23-case behavior eval** firing curated prompts at the live model and asserting in-domain answering, off-domain refusal, and medical-advice refusal: **23/23 passing**; backed by 80 unit tests.
- Enforced **per-user data isolation**: every tool resolves the patient from the server-derived session, never a client-supplied ID, so no prompt reaches another patient's protected health information.

Projects

Gujarati LM Suite: GPT-2 + Generative QA Base Model | Fine-tuned
Trained a 52M-parameter GPT-2 causal language model from scratch for Gujarati with a custom 64K-vocab SentencePiece tokenizer; fine-tuned it on IndicSQuAD for generative question answering using a structured `Context / Question / Answer` template, cutting validation loss 18% over 3 epochs. *Tech: PyTorch, HuggingFace Transformers, SentencePiece, FP16*

Gujarati LLaMA2 Base HuggingFace
Designed a compact LLaMA2-architecture language model (6 layers, 512 hidden dim, 64K vocab) trained on ∼2GB of Gujarati text; implemented custom learning-rate scheduling with gradient-norm tracking. *Tech: PyTorch, HuggingFace*

Mars-Mars Game RL GitHub
Built a 2D platformer in Pygame wrapped as a Gymnasium-compatible reinforcement-learning environment; trained a Proximal Policy Optimization (PPO) agent with shaped rewards handling dynamic wind, jetpack physics, and fuel management. *Tech: Python, Pygame, Gymnasium, Stable-Baselines3*

PromptCraft Live
Deployed web tool powered by the Google Gemini API generating tailored coding prompts; 295+ organic users over 14 months with zero paid marketing. *Tech: Gemini API, JavaScript, TailwindCSS*

Skills

AI & ML: LLM tool calling / function calling, agent loops, prompt engineering, prompt guardrails, LLM behavior evaluation, LLM fine-tuning, PyTorch, HuggingFace Transformers, SentencePiece, Reinforcement Learning, Scikit-learn

Languages: Python, TypeScript, JavaScript, C/C++, SQL

Frameworks: Fastify, Flask, React Native, AngularJS, TailwindCSS

Cloud & DevOps: Azure OpenAI, Azure DevOps (Boards, Repos, Pipelines), Azure Container Apps, Application Insights, PostHog, Docker, Git

Tools & Databases: Postman, Cursor, Figma, VS Code, PostgreSQL, MySQL, MongoDB

Education

B.Tech. Computer Engineering (Lateral Entry), Charusat University 2024 - 2027 (Expected)
Diploma in Computer Engineering, B & B Institute of Technology 2021 - 2024